

Ideation Phase

Brainstorm & Idea Prioritization Template

Date	31 January 2025
Team ID	
Project Name	AI FAQ Assistant API
Maximum Marks	4 Marks

Brainstorm & Idea Prioritization Template:

Brainstorming gives the team a free and open space to share ideas for AI FAQ Assistant API, a NestJS backend that stores a business's FAQs and uses AI to answer customer questions accurately from them. Volume is valued over polish; unusual ideas are welcome and every idea can be built upon.

Reference: <https://www.mural.co/templates/brainstorm-and-idea-prioritization>

Step-1: Team Gathering, Collaboration and Select the Problem Statement

Selected problem statement: Support teams answer the same customer questions again and again, customers wait for replies, and there is no simple API that turns a company's own FAQs into accurate, trustworthy AI answers.

- The team gathers (in person or remotely) and agrees on ground rules: one idea per note, no criticism during ideation, build on others' ideas.
- Roles: Facilitator, Timekeeper, Scribe, Backend (NestJS) lead, AI/prompt lead.

Step-2: Brainstorm, Idea Listing and Grouping

Group	Ideas
FAQ management	Create, read, update, delete FAQs Categories and tags Bulk import from CSV / JSON Draft / Published / Archived status Version history of answers
AI answering	Semantic search with embeddings (matches meaning, not just words) Short answer written from the top matching FAQs, with source links Confidence score on every answer "I don't know" fallback instead of guessing Follow-up questions using session context Answers in the customer's language
Feedback & escalation	Thumbs up / down on each answer Escalate to a human agent by ticket or email Queue of unanswered questions Suggest new FAQs from unanswered questions
Analytics	Top asked questions Unanswered rate and satisfaction score Usage per API key

Group	Ideas
Security & platform	JWT login for admins and editors API keys per organisation Rate limiting and Redis answer cache Swagger / OpenAPI documentation PostgreSQL with pgvector, Docker, unit and e2e tests

Step-3: Idea Prioritization

Idea	Impact	Effort	Priority
FAQ CRUD with status workflow	High	Low	P1 – Must have (MVP)
JWT auth and roles	High	Medium	P1 – Must have (MVP)
Embeddings and semantic FAQ search	High	Medium	P1 – Must have (MVP)
AI answer grounded in FAQs, with sources	High	Medium	P1 – Must have (MVP)
Confidence score and “I don’t know” fallback	High	Medium	P1 – Must have (MVP)
Swagger API documentation	Medium	Low	P1 – Must have (MVP)
API keys per organisation	High	Medium	P2 – Should have
Follow-up questions (session context)	Medium	Medium	P2 – Should have
Feedback and human escalation	Medium	Medium	P2 – Should have
Redis cache and rate limiting	Medium	Medium	P2 – Should have
Bulk import from CSV / JSON	Medium	Low	P2 – Should have
Analytics (top and unanswered questions)	Medium	Medium	P3 – Later
Multi-language answers	Low	High	P3 – Later
Suggest new FAQs from unanswered questions	Low	High	P3 – Later